

Log to Incident

Table of Contents

1. Purpose	1
2. Log to Incident	2
2.1. What the process does	2
2.2. What must be agreed up front	2
2.3. Client and server applications	3
2.4. Platform implementations	3
3. Using ServiceNow and Dynatrace	4
3.1. End-to-end flow	4
3.1.1. Step-by-step configuration	5
3.1.2. Client-Server Complications	6
3.2. Step 0 — ServiceNow Modeling	7
3.2.1. ServiceNow Modeling	8
3.2.2. Snippets from the CSDM-CMDB model	10
3.2.3. ServiceNow — CSDM Application Services	10
3.2.4. ServiceNow — Kubernetes pod CIs and tags	11
3.2.5. CMDB model	12
3.2.6. ServiceNow — SGC topology import and sys_object_source	12
3.2.7. Dynatrace — management zone	13
3.2.8. Step 0 completion criteria	13
3.3. Step 1 — Application emits a log line with requisite metadata	14
3.3.1. Required log content	14
3.3.2. Log4j configuration (reference)	14
3.3.3. Log emission	15
3.3.4. Log path and OneAgent tail (not NFS)	16
3.3.5. Dynatrace and ServiceNow (Step 1)	16
3.4. Step 2 — Dynatrace detects ERROR and WARN log lines	17
3.4.1. OpenPipeline processor flow	17
3.4.2. Dynatrace objects to configure	17
3.4.3. Custom log source	17
3.4.4. Logs routing	18
3.4.5. Log alerts pipeline (complete reference)	18
3.4.6. Davis processor properties	20
3.4.7. Example Grail log record at Davis input	20
3.4.8. Alerting profile	20
3.4.9. Host CPU metric event (optional companion path)	21
3.4.10. ServiceNow (Step 2)	22
3.5. Step 3 — Dynatrace sends the problem to ServiceNow	22

3.5.1. Webhook endpoint	22
3.5.2. Problem notification payload template	22
3.5.3. Webhook field mapping	22
3.5.4. Complete example webhook body	23
3.5.5. Ensuring log fields reach ServiceNow	24
3.5.6. ServiceNow (Step 3)	24
3.6. Step 4 — ServiceNow receives the event and binds infrastructure CI	25
3.6.1. Event listener	25
3.6.2. em_event field mapping	25
3.6.3. CI binding procedure	26
3.6.4. Entity type to CMDB class	26
3.6.5. Description and alert text	26
3.6.6. SGC and Event Management configuration	27
3.6.7. Application and Dynatrace (Step 4)	27
3.7. Step 5 — Create incident bound to Application Service	27
3.7.1. Severity policy	27
3.7.2. Layer model	28
3.7.3. Incident correlation	28
3.7.4. CSDM identifier lookup	29
3.7.5. Known limitation — Davis problem bundling	30
3.7.6. ServiceNow incident automation specification	30
3.7.7. Business rule order and rationale	31
3.7.8. Alert to incident linkage (em_alert.incident)	31
3.7.9. Problem and alert lifecycle (OPEN / RESOLVED)	31
3.7.10. OpenPipeline and custom-source attributes (generic Log4j pattern)	32
3.7.11. Event Management rules (manual configuration)	32
3.7.12. Script Includes	33
3.7.13. Business rules (prescriptive configuration)	34
3.7.14. Ansible deployment	36
3.7.15. Example resolution (Kubernetes worker pod)	36
3.7.16. Application and Dynatrace (Step 5)	37
Appendix A: Legacy and SGO-Dynatrace coexistence	38
A.1. Separation by source	38
A.2. Scoping rules and profiles	38
A.3. Host CPU alerts (not Spark-specific)	38
A.4. When to retire legacy connector	39

Chapter 1. Purpose

Large incident resolution times hurt business operations and customer satisfaction. One way to shorten resolution is to bring critical application problems to the right operations team quickly and with enough context to act. That requires bridging **observability** (detecting and describing problems from runtime signals) with **IT service management** (tracking work as incidents, assignment, and closure).

This document defines that bridge as the **Log to Incident** process: selected application log messages become managed incidents for the application operations team. The process is named for the **signal type** and **outcome**, not for any particular vendor or tool. Implementations may use different observability and ITSM platforms; this document's worked example uses Apache Spark as the application, with platform-specific sections for how those platforms are configured.

Chapter 2. Log to Incident

Log to Incident is the Observability process that detects and turns **critical** application log messages into **tracked** IT incidents so the application operations team can respond in time. It does this through a four phase process:

1. The application logs a message,
2. The observability tooling detects the message as a being salient and creates an alert
3. The alert is transmitted to the ITSM platform
4. The ITSM platform determines which alerts are salient and then creates an incident for operators to track.

2.1. What the process does

1. An application writes operational messages to its application log (errors, warnings, or other agreed severity).
2. Observability tooling collects those messages and decides which ones are severe enough to treat as a **problem** (a detected issue worth escalating).
3. The problem is handed to IT service management as an **event** or **alert** tied to the right infrastructure or application identity.
4. Automation creates (or updates) an **incident** assigned to the application's support organization, carrying enough log context for diagnosis.
5. The operations team works and closes the incident through normal ITSM practice.

The process intentionally does **not** depend on a single observability product or a single ITSM product. What must stay stable is the **contract**: which log lines matter, how the application is identified for assignment, and what context travels with the incident.

2.2. What must be agreed up front

Not every log line should become an incident. Flooding operations with low-value tickets increases noise and lengthens real outages. Observability engineers **must** agree with application developers and operations on:

- Which severities or message patterns are **incident-worthy**.
- How log lines are structured so the incident carries actionable context (what failed, where, which workload or component).
- Which **application service** (or equivalent business CI) owns the work when an incident is created.

2.3. Client and server applications

Many systems have both **service-side** components (clustered services, pods, long-running servers) and **client-side** components (drivers, desktops, mobile apps, or other unmanaged endpoints). Log to Incident must still escalate both, even when the client is not a managed CI in the CMDB. That usually means:

- Service-side: map from the running workload's identity to the application service that operations supports.
- Client-side: map from a durable log-path or naming contract to a logical application service that represents the client population.

2.4. Platform implementations

Later chapters describe how to realize Log to Incident on specific platforms. The first implementation in this document is **Using ServiceNow and Dynatrace**. Additional chapters (for example Elastic Stack and ServiceNow) **may** be added without renaming the process.

Chapter 3. Using ServiceNow and Dynatrace

This chapter is the reference implementation of Log to Incident using **Dynatrace** as the observability platform and **ServiceNow** as the ITSM / Event Management platform. Apache Spark is the sample application (WARN and ERROR log lines escalate in the lab; production agreements often limit escalation to ERROR after review with operations and development).

External references (vendor documentation):

- [Dynatrace — Log Analytics and OpenPipeline](#)
- [Davis problem detection](#)
- [ServiceNow — Event Management](#)
- [Apache Spark Documentation](#)

3.1. End-to-end flow

The figure below illustrates the end-to-end flow of this implementation: the application JVM writes to the application log file; Dynatrace detects, correlates, and notifies; ServiceNow creates the event, alert, and incident.




Figure 1. Log to Incident end-to-end data flow

Critical to the process is selecting the best configuration items (CIs) to assign to the alert and even more importantly the incident. CI selection is especially important for the incident as the incident will be directed to the support group of the CI. Care must be taken to make sure that the CI then represents the application and the support group on the application CI the application's operations team.

3.1.1. Step-by-step configuration

The end-to-end data flow in the figure above maps into six major areas of configuration required to implement the Log to Incident flow. The table below defines the the six majors steps required to configure the Log to Incident flow.


Step	Layer	Description
0	ServiceNow Modeling	The CSDM-CMDB model provides the foundation that allows a data-driven approach to configure the Log to Incident flow independent of any particular application being observed. It consists of the ServiceNow service and application CIs used to represent the application, the CIs injected into the CMDB by Dynatrace the workload CIs, and those CIs discovered by the Kubernetes cluster used to run the Spark platform, <code>sys_object_source</code> ; pattern-specific CMDB model (CMDB model)
1	Application / Log4j	Spark utilized Log4j for logging. These log files must be configured to be emitted in a way compatible Dynatrace's expectation of where and how to expect the logs.
2	Dynatrace detection	Configuration of the OpenPipeline that matches only the WARN or ERROR log messages.
3	Dynatrace → ServiceNow	Configuration of the Webhook that carries the log line and impacted entity to ServiceNow as well as the layout of the event record.
4	ServiceNow Event Management	Specification of the <code>em_event</code> / <code>em_alert</code> and how to bind the correct infrastructure CI to the alert.
5	ServiceNow ITSM	Configuration of ServiceNow to automatically create an incident for the alert and to determine the right CI to bind to the incident.

3.1.2. Client-Server Complications

You can deploy Spark as a client-server application or a server-only application. We had chosen enable the client-server mode to allow data scientists or engineers to launch applications from their own laptops or workstations. This allows for a more flexible and efficient development and testing environment. It also provides a more robust observability exercise as we have to address the client use case, which can occur in many other applications. For example, you can see the client in mobile apps where the client is not only a separate application, but it is an applications that runs on a device that is not managed in ServiceNow's CMDB.

The end-to-end flow in Figure XXX must handle both the **Service-side** and the **Client-side** use cases. The CSDM model for the application will differ between a client and the application service. Where the log files reside will be different between a client application and the application service. Lastly, the way in which the right CI is mapped to an incident will differ, too. As such, Steps 0, 1, and 5 must handle both the Service-side and the Client-side use cases.

For those not familiar with Spark's client-server model, the figure below illustrates the different components of the Spark application.



Spark can run in one of two modes: **Client Mode** or **Cluster Mode**. In client mode, the client-side application is a standalone application that runs on a workstation or host. The client then communicates with the Spark Master running as a service and the executors running on the Spark Workers. The Spark Master and the workers run on the service side. The executors on the workers run jobs and constitute the heart of the application.

In cluster mode, the Spark Driver runs on a Kubernetes pod and communicates with the Spark Master running as a service and the executors running on the Spark Workers. The Spark Master and the workers run on the service side. The executors on the workers run jobs and constitute the heart of the application.

Clean up the sections below:

Typically, the client-side application is not managed by ServiceNow's CMDB. Also, it does not run on a Kubernetes pod. Therefore, the CI that is mapped to the incident will be a manual logical service that is defined in the CSDM model. The log file will be located in a fixed path segment on the host. The way in which the right CI is mapped to an incident will be done by the application developer by setting the `incident.cmdb_ci` field on the alert and incident records.

Both patterns require the fully formatted log line in event, alert, and incident descriptions; Dynatrace **ERROR** or **WARN** detection; and `incident.cmdb_ci = cmdb_ci_service_discovered` (Application Service), not the infrastructure CI on the alert.

Both patterns follow the same observability path from application log line to ServiceNow incident. Figure 1 is the navigation map; pattern-specific configuration is called out in Steps 0, 1, and 5.

Steps 2–4 are shared; Steps 0, 1, and 5 include dual **Service-side** and **Client-side** subsections.

3.2. Step 0 — ServiceNow Modeling

Before log problems can bind to infrastructure CIs and incidents to an **Application Service**, we must model the application in ServiceNow. This involves defining the CSDM hierarchy and discovering pattern-specific CMDB components. We need to do this for the Service-side and the


Client-side components.

3.2.1. ServiceNow Modeling

The ServiceNow modeling is the data foundation of the Log to Incident process. It consists of the CSDM hierarchy and the pattern-specific CMDB components. We need to do this for the Service-side and the Client-side components.

3.2.1.1. Service-side

The Service-side model shows logical relationships between CSDM Application Services, Kubernetes-discovered CIs, SGC-imported entities, and Event Management bindings for pod-scoped logs. Use it to verify **em_alert.cmdb_ci** → pod and **incident.cmdb_ci** → Application Service before enabling automation.



Lab scenario: WARN from **Master:** logger in `/mnt/spark/logs/spark-master-0/spark-app.log`. ServiceNow binds **em_event.cmdb_ci** / **em_alert.cmdb_ci** to **cmdb_ci_kubernetes_pod**. **incident.cmdb_ci** is **cmdb_ci_service_discovered** for Application Service **spark-master** (from `servicenow/regions/brooks-lab/spark.csdm.yaml`), via the **Spark Service-Side Log Pattern** — not the Dynatrace problem object and not the pod CI.

The table below lists each CMDB object in the Service-side scenario and which Event Management / ITSM field it populates.

Object	Role in this scenario
cmdb_ci_business_app — Data and Analytic Services	CSDM parent; not an incident CI
cmdb_ci_service — Apache Spark	Business Service parent of Application Service
cmdb_ci_service_discovered — spark-master	incident.cmdb_ci (preferred)

CSDM fragment (brooks-lab):

```
- name: Spark Client
  identifier: spark-client
  platform: host
  service_mapping: manual
  discover: false
  depends_on:
    - Spark Master
```

3.2.2. Snippets from the CSDM-CMDB model

The step table maps each layer to the section that configures it. Steps 2–4 are shared; Steps 0, 1, and 5 include parallel **Service-side** and **Client-side** subsections.

Log to Incident supports **two incident-mapping patterns** that share the same Dynatrace OpenPipeline → Davis → ServiceNow Event Management pipeline (Steps 2–4) but diverge at **CMDB prerequisites** (Step 0), **log emission** (Step 1), and **incident correlation** (Step 5):

- **Spark Service-Side Log Pattern** — Kubernetes pod logs; **incident.cmdb_ci** via **Contains** traversal from pod CI to Application Service.
- **Spark Client-Side Log Pattern** — PySpark driver logs on any host; **incident.cmdb_ci** via **log path contract** to a manual logical Application Service.

Use the comparison table below to see where each pattern applies and how incidents resolve. Detailed CMDB models and correlation algorithms appear in [CMDB model](#) and [Incident correlation](#).

3.2.3. ServiceNow — CSDM Application Services

Declare one Application Service per logical service boundary. Each Spark Kubernetes workload that should roll up to that service carries a matching pod label synced to **cmdb_key_value**.

The table maps brooks-lab Spark Application Service identifiers to pod roles and the label each pod must carry.

Application Service identifier	Workload role	Tag on pod (label key=value)
spark-master	Spark master StatefulSet	servicenow.io/application-service-identifier=spark-master
spark-history-server	Spark history server	servicenow.io/application-service-identifier=spark-history-server
spark-worker-lab1	Spark worker on Lab1	servicenow.io/application-service-identifier=spark-worker-lab1
spark-worker-lab2	Spark worker on Lab2	servicenow.io/application-service-identifier=spark-worker-lab2

Reference CSDM fragment from [servicenow/regions/brooks-lab/spark.csdm.yaml](#):

```

application_services:
- name: Spark Master
  identifier: spark-master
  short_description: Spark master – Kubernetes StatefulSet
  operational_status: "1"
  parent_business_service: Apache Spark
  platform: kubernetes
  service_mapping: tags
  discover: false
  service_tier: app
  environment: on-prem
  location: brooks-lab
  cluster: brooks-lab
  namespace: spark
  owned_by: gbrooks

- name: Spark Worker (Lab1)
  identifier: spark-worker-lab1
  short_description: Spark worker on Lab1
  platform: kubernetes
  service_mapping: tags
  discover: false
  namespace: spark

```

CSDM deployment must create:

- cmdb_ci_service_discovered rows and Business Application → Business Service → Application Service **Contains** hierarchy
- Declared **depends_on** edges (shared infrastructure the service depends on)
- tag_list on each tag-based Application Service (Service Mapping Operations REST) so the mapping engine knows which cmdb_key_value labels qualify

3.2.4. ServiceNow — Kubernetes pod CIs and tags

Pod manifest labels are the single source of truth for tag-based Service Mapping on Spark workloads.

Activity	Produces
Kubernetes discovery / KVA resync	cmdb_ci_kubernetes_pod rows for running pods
Label sync to CMDB (REST upsert on pod CI)	cmdb_key_value rows for servicenow.io/* labels copied from pod metadata.labels

Pod manifest labels (brooks-lab Spark worker example — apply at deploy time):

```

apiVersion: v1
kind: Pod
metadata:
  name: spark-worker-lab1-0
  labels:
    servicenow.io/application-service-identifier: spark-worker-lab1
    servicenow.io/environment: on-prem
    servicenow.io/location: brooks-lab
    app.kubernetes.io/name: spark-worker
spec:
  containers:
    - name: spark-worker

```

```

env:
  - name: POD_NAME
    valueFrom:
      fieldRef:
        fieldPath: metadata.name
volumeMounts:
  - name: app-logs
    mountPath: /opt/spark/logs
    subPathExpr: ${POD_NAME}
volumes:
  - name: app-logs
    persistentVolumeClaim:
      claimName: spark-logs-pvc

```

After tags exist on pod CIs, run **tag-based Service Mapping** (scheduled job or **Update map** in Service Mapping Workspace) to materialize:

cmdb_ci_service_discovered → **Contains::Contained by** → cmdb_ci_kubernetes_pod

3.2.5. CMDB model

Alerts bind to **infrastructure** CIs; incidents bind to **Application Services**. Each Log to Incident pattern requires a different CMDB shape so operators assign the most informative CI at each layer. The subsections below define **Service-side** and **Client-side** models side by side.

3.2.6. ServiceNow — SGC topology import and sys_object_source

SGC scheduled import	Dynatrace entity type	CMDB class	Role in log incidents
Hosts	HOST	cmdb_ci_linux_server	Fallback when log tails on host and pod entity unavailable
Kubernetes Pod	CLOUD_APPLICATION_INSTANCE	cmdb_ci_kubernetes_pod	Preferred alert/event CI for pod-scoped logs
Process Groups	PROCESS_GROUP	cmdb_ci_appl	Alternative when Davis binds to JVM process
Application Relationships	Smartscape edges	cmdb_rel_ci	Optional topology context

Each imported entity receives a **sys_object_source** row: name=SGO-Dynatrace, id=<Dynatrace entityId>, target_sys_id → CMDB CI sys_id.

Configure on the ServiceNow instance:

- Service Graph Connector for Observability – Dynatrace (scoped app sn_dynatrace_integ)
- Inbound event listener: POST /api/sn_em_connector/em/inbound_event?source=SGO-Dynatrace

- Scheduled imports for hosts, Kubernetes pods, and process groups covering the cluster namespace

3.2.7. Dynatrace — management zone

Place hosts, Kubernetes cluster entities, and pods that run the application in a **management zone** dedicated to the integration (for example **Application Observability**). The **alerting profile** for ServiceNow problem notifications must scope to that management zone so only in-scope problems are forwarded.

3.2.8. Step 0 completion criteria

Complete the shared criteria first, then the pattern-specific rows for each Log to Incident pattern you deploy.

#	Criterion (shared)
1	SGC imported hosts, Kubernetes pods, and process groups; <code>sys_object_source</code> rows exist for entity IDs in the management zone
2	Dynatrace management zone scopes hosts and pods that emit qualifying logs
3	Event Management rules mark SGO-Dynatrace log events Ready when <code>cmdb_ci</code> is populated (Event Management rules)
4	Incident business rules deployed: pod bind (K8s log pod bind business rules), create incident (Incident correlation)

3.2.8.1. Service-side

#	Criterion
1	Application Service CIs exist for each K8s Spark role; <code>tag_list</code> configured for tag-based mapping
2	Pod labels synced to <code>cmdb_key_value</code> on <code>cmdb_ci_kubernetes_pod</code>
3	Service Mapping materialized Contains from Application Service to member pods
4	Incident automation implements Service-Side Contains traversal (Service-side pattern)

3.2.8.2. Client-side

#	Criterion
1	Application Service Spark Client exists with <code>service_mapping: manual</code> and <code>service_status=operational</code>
2	Client drivers write under <code>/mnt/spark/logs/spark-client/<instance>/</code> (Client-side log emission)
3	Dynatrace custom log source and OpenPipeline list spark-client paths before generic <code>/mnt/spark/logs/*/</code>
4	Incident automation implements <code>resolveFromSparkClientLogPath()</code> (Client-side pattern)

3.3. Step 1 — Application emits a log line with requisite metadata

3.3.1. Required log content

Attribute	Source	Consumed by
Timestamp, level, logger, line number, message	Log4j PatternLayout (or equivalent)	Dynatrace content field; ServiceNow em_event.description
Workload directory in file path	Kubernetes subPathExpr: \$(POD_NAME) on log volume	OpenPipeline path parser → pod name
Shared storage path	PVC or hostPath mount visible to OneAgent	Dynatrace custom log source pattern

3.3.2. Log4j configuration (reference)

Use a **RollingFile** appender only (no console appender in production pods). Set **APP_LOG_DIR** or **SPARK_LOG_DIR** to the pod-specific directory.

Key fields: **fileName** and **filePattern** path prefixes **must** match the Dynatrace custom log source paths in [custom log source specification](#) and the OpenPipeline **log.source** matchers in [OpenPipeline specification](#) (brooks-lab: `/mnt/spark/logs/<pod>/spark-app.log` and rotated `spark-app-*.log`).

Specification 2.4.2-1. Log4j RollingFile (Spark client / chapter reference)

```
# RollingFileAppender – file logging
appender.rolling.type = RollingFile
appender.rolling.name = RollingFile
appender.rolling.fileName = ${env:SPARK_LOG_DIR:-/mnt/spark/logs/client}/spark-app.log
appender.rolling.filePattern = ${env:SPARK_LOG_DIR:-/mnt/spark/logs/client}/spark-app-%d{yyyy-MM-dd}-%i.log
appender.rolling.layout.type = PatternLayout
appender.rolling.layout.pattern = %d{yyyy-MM-dd HH:mm:ss} %-5p %c{1}:%L - %m%n%ex
```

Pass the configuration to the JVM:

```
-Dlog4j2.configurationFile=file:///opt/myapp/conf/log4j2.properties
```

Optional: add Mapped Diagnostic Context (MDC) for pod name in the log line itself (redundant when the path encodes pod identity, but useful if log.source is normalized to a wildcard pattern):

```
appender.rolling.layout.pattern = %d{yyyy-MM-dd HH:mm:ss} %-5p [%X{podName}] %c{1}:%L - %m%n%ex
```

Set MDC at process startup from the `POD_NAME` environment variable.

3.3.3. Log emission

Log emission defines where application JVMs write ERROR/WARN lines so OneAgent can tail them and OpenPipeline can extract workload identity. Requirements diverge by pattern; configure **Service-side** and **Client-side** subsections as applicable.

3.3.3.1. Service-side

Service-side workloads **must** run as **Kubernetes pods** with pod-scoped log directories visible on the node. The requirement table lists each constraint and why it matters for path-based pod binding.

Requirement	Specification
Runtime	Workload runs as a Kubernetes pod (not a host-local client driver)
Pod identity	<code>metadata.name</code> is the directory name in the log file path
Log4j appender	RollingFile only in production pods; no console-only logging for incident-bound workloads
Log directory	<code>subPathExpr: \$(POD_NAME)</code> on a volume mounted at a path OneAgent can tail on the node (brooks-lab: <code>/mnt/spark/logs/\$(POD_NAME)/</code>)
Log file name	<code>application.log</code> or rotated <code>application-.log</code> (brooks-lab Spark: <code>spark-app.log</code> , <code>spark-app-.log</code>)
Pod label	<code>servicenow.io/application-service-identifier=<identifier></code> matching a CSDM Application Service (Step 0)
Levels	Log4j emits ERROR and WARN for conditions that should open incidents per severity policy

The init container or `subPathExpr` mount must ensure each pod writes to a unique directory whose name equals `metadata.name`:

```
env:
- name: APP_LOG_DIR
  value: /var/log/myapp/$(POD_NAME)
- name: POD_NAME
  valueFrom:
    fieldRef:
      fieldPath: metadata.name
volumeMounts:
- name: app-logs
  mountPath: /var/log/myapp
  subPathExpr: $(POD_NAME)
```

Resulting path on the node (OneAgent tail source):

```
/mnt/spark/logs/myapp-worker-7d4f9b-xk2lm/application.log
└─ must equal cmdb_ci_kubernetes_pod.name
```

3.3.3.2. Client-side

The **Spark Client-Side Log Pattern** requires a fixed path segment (**spark-client**) and one subdirectory per driver JVM. The attribute table maps each path element to the consumer that resolves incidents.

Attribute	Source	Consumed by
Timestamp, level, logger, message	Log4j PatternLayout	Dynatrace loglevel , content ; ServiceNow em_event.description
Fixed AS path segment spark-client	run-chapters.sh sets SPARK_LOG_DIR=/mnt/spark/logs/spark-client/<instance>	OpenPipeline Davis matcher; incident BR (Client-side pattern)
Per-driver instance directory	{hostname}-{pid} (override with SPARK_DRIVER_INSTANCE)	OneAgent enrichment → spark.client.instance
Optional process tag	DT_TAGS=servicenow.io/application-service-identifier=spark-client	Process-group problems only; not required for log-path incidents

Example layout:

```
/mnt/spark/logs/spark-client/lab3-12345/spark-app.log  
/mnt/spark/logs/spark-client/lab3-12345/spark-app-2026-07-06-1.log
```

Do not colocate multiple concurrent drivers in one directory — Log4j RollingFile would interleave or truncate **spark-app.log**.

3.3.4. Log path and OneAgent tail (not NFS)

Dynatrace **does not read logs from the NFS server**. OneAgent on each **Kubernetes node** tails log files on the **node's local filesystem**:

1. Inside the pod, Log4j writes to a pod-local mount (for example **/opt/spark/logs/\$(POD_NAME)/spark-app.log**).
2. A **hostPath** (or equivalent node-visible mount) maps that directory to **/mnt/spark/logs/\$(POD_NAME)/** on the node where the pod is scheduled.
3. OneAgent **custom log source** **LOG_PATH_PATTERN** matches **/mnt/spark/logs/*/...** on the **host** and tails those files into Grail.
4. NFS volumes used for shared Spark data (**/mnt/spark/events**, **/mnt/spark/data**, ...) are **not** the Dynatrace log tail source.

3.3.5. Dynatrace and ServiceNow (Step 1)

No Dynatrace or ServiceNow configuration at log emission time. **Preconditions**: OneAgent can read the log path on the node; the path directory name equals **cmdb_ci_kubernetes_pod.name** after KVA discovery; pod carries **servicenow.io/application-service-identifier**; tag-based Service Mapping has materialized **Contains::Contained by** from Application Service to pod **before** the first qualifying log line (Step 0).

3.4. Step 2 — Dynatrace detects ERROR and WARN log lines

Logs ingested through **OpenPipeline** are evaluated by pipeline processors. Classic builtin:logmonitoring.log-events rules do **not** apply to logs routed through OpenPipeline.

3.4.1. OpenPipeline processor flow



Figure 2. OpenPipeline log alerts processor chain

3.4.2. Dynatrace objects to configure

Object	Settings schema	Purpose
Custom log source	builtin:logmonitoring.custom-log-source-settings	OneAgent tails application log paths on hosts or nodes
Logs routing	builtin:openpipeline.logs.routing	Route matched logs into the log-alerts pipeline
Log alerts pipeline	builtin:openpipeline.logs.pipelines	Extract pod name, set dt.source_entity, Davis processor
Alerting profile	builtin:alerting.profile	Forward in-scope problems to ServiceNow notification
Problem notification	builtin:problem.notifications	POST webhook on problem OPEN and RESOLVED

3.4.3. Custom log source

Register path patterns that cover per-workload directories and rotated files.

Key fields: **path** patterns **must** align with Log4j **fileName** / **filePattern** (Log4j RollingFile specification) and OpenPipeline **log.source** matchers (OpenPipeline specification). **Pod** paths use segment = `cmdb_ci_kubernetes_pod.name`. **Client-mode** paths use `/mnt/spark/logs/spark-client/<instance>/` (Client-side pattern) with enrichment `spark.client.as_identifier` and `spark.client.instance`. OpenPipeline **drops** legacy `-chapter` paths parsed as pod names. **enrichment** attributes `k8s.log.path` (`{0}`) and `k8s.pod_name` or `spark.client.instance` (`{1}`) **must** be set on wildcard paths – list **more specific** patterns (spark-client) **before** generic `/mnt/spark/logs//`.

Specification 2.5.1-1. Custom log source (brooks-lab Spark)

```
{
  "schemaId": "builtin:logmonitoring.custom-log-source-settings",
  "scope": "environment",
```

```

"value": {
  "enabled": true,
  "config-item-title": "Spark Lab - application log files",
  "custom-log-source": {
    "type": "LOG_PATH_PATTERN",
    "values-and-enrichment": [
      {
        "path": "/mnt/spark/logs/*/spark-app.log",
        "enrichment": [
          { "type": "attribute", "key": "k8s.log.path", "value": "${0}" },
          { "type": "attribute", "key": "k8s.pod_name", "value": "${1}" }
        ]
      },
      {
        "path": "/mnt/spark/logs/*/spark-app-*.log",
        "enrichment": [
          { "type": "attribute", "key": "k8s.log.path", "value": "${0}" },
          { "type": "attribute", "key": "k8s.pod_name", "value": "${1}" }
        ]
      },
      { "path": "/opt/spark/logs/spark-app.log*" }
    ]
  }
}

```

3.4.4. Logs routing

Route OneAgent-sourced records whose log.source matches the application path into the alerts pipeline:

```

{
  "schemaId": "builtin:openpipeline.logs.routing",
  "scope": "environment",
  "value": {
    "routing": [
      {
        "enabled": true,
        "description": "Application ERROR/WARN log alerts",
        "matcher": "matchesValue(log.source, \"\n/var/log/myapp/*\")",
        "pipelines": ["application-log-alerts"]
      }
    ]
  }
}

```

3.4.5. Log alerts pipeline (complete reference)

The pipeline has three processor stages before storage:

1. **Fields extract** — parse `k8s.pod_name` from `log.source` when the path is literal.
2. **Fields add** / **entity lookup** — set `dt.source_entity` via `lookupEntity(CLOUD_APPLICATION_INSTANCE, k8s.pod_name)`.
3. **Davis processor** — match **WARN/ERROR**; set `event.name` to include pod name; emit `event.description` with full log line.
4. **Bucket assignment** — store in `default_logs` bucket.

Key fields: DQL **matcher** on **log.source**; **lookupEntity**; Davis **event.name**, **event.description**, **dt.source_entity**.

Specification 2.5.4-1. OpenPipeline log alerts pipeline (brooks-lab Spark)

```
{
  "schemaId": "builtin:openpipeline.logs.pipelines",
  "scope": "environment",
  "value": {
    "displayName": "Spark Lab - log alerts",
    "customId": "spark-lab-log-alerts",
    "processing": {
      "processors": [
        {
          "id": "extract-spark-pod-from-path",
          "type": "dql",
          "matcher": "matchesValue(log.source, \"/mnt/spark/logs/*\") OR matchesValue(log.source, \"/opt/spark/logs/*\")",
          "dql": {
            "script": "parse log.source, \"LD '/mnt/spark/logs/' LD:k8s.pod_name LD '/spark-app.log'\""
          }
        },
        {
          "id": "resolve-spark-pod-entity",
          "type": "fieldsAdd",
          "matcher": "isNotNull(k8s.pod_name)",
          "fieldsAdd": {
            "fields": [
              {
                "name": "dt.source_entity",
                "value": "lookupEntity(CLOUD_APPLICATION_INSTANCE, k8s.pod_name)"
              }
            ]
          }
        }
      ]
    },
    "davis": {
      "processors": [
        {
          "matcher": "(loglevel == \"WARN\" OR loglevel == \"ERROR\") AND (matchesValue(log.source, \"/mnt/spark/logs/*\") OR matchesValue(log.source, \"/opt/spark/logs/*\"))",
          "davis": {
            "properties": [
              { "key": "event.type", "value": "ERROR_EVENT" },
              { "key": "event.name", "value": "Spark log {loglevel} on {k8s.pod_name}" },
              { "key": "event.description", "value": "Spark {loglevel} on {log.source}: {content}" },
              { "key": "dt.source_entity", "value": "{dt.source_entity}" }
            ]
          }
        }
      ]
    }
  }
}
```



OpenPipeline **Add fields (fieldsAdd)** accepts **lookupEntity(CLOUD_APPLICATION_INSTANCE, <pod-name-field>)** in the API **value** property (Settings API uses **name**, not **field**, for the target attribute). DQL processors in the Processing stage do **not** enable **fetch** or **smartscapeNodes**; do not use Grail notebook lookup syntax there. When lookup misses, **dt.source_entity** remains the OneAgent HOST. The required outcome is **dt.source_entity =**

`CLOUD_APPLICATION_INSTANCE-...` before the Davis processor runs.

3.4.6. Davis processor properties

Property	Value
Matcher	(loglevel == "WARN" OR loglevel == "ERROR") AND matchesValue(log.source, "/var/log/myapp/*")
event.type	ERROR_EVENT
event.name	Spark log {loglevel} on {k8s.pod_name}
event.description	Application {loglevel} on {log.source}: {content}
dt.source_entity	CLOUD_APPLICATION_INSTANCE entity when pod resolved; else pass-through process or host
dt.davis.event_timeout	15m (adjust to problem correlation needs)

3.4.7. Example Grail log record at Davis input

Attribute	Example
content	2026-06-27 14:53:46 ERROR Worker:264 - All masters are unresponsive! Giving up.
loglevel	ERROR
log.source	/var/log/myapp/myapp-worker-7d4f9b-xk2lm/application.log
workload.pod_name	myapp-worker-7d4f9b-xk2lm (after fields extract)
dt.source_entity	CLOUD_APPLICATION_INSTANCE-A1B2C3D4E5F67890
host.name	k8s-node-01 (tail host — not the alert CI when pod entity is set)

3.4.8. Alerting profile

Scope the profile to the integration management zone and include ERROR and CUSTOM_ALERT severity levels (Davis log events map to ERROR):

```
{
  "schemaId": "builtin:alerting.profile",
  "scope": "environment",
  "value": {
    "name": "Application Observability - ServiceNow",
    "severityRules": [
      { "severityLevel": "AVAILABILITY", "delayInMinutes": 0, "tagFilterIncludeMode": "NONE" },
      { "severityLevel": "ERRORS", "delayInMinutes": 0, "tagFilterIncludeMode": "NONE" },
      { "severityLevel": "PERFORMANCE", "delayInMinutes": 0, "tagFilterIncludeMode": "NONE" },
      { "severityLevel": "RESOURCE_CONTENTION", "delayInMinutes": 0, "tagFilterIncludeMode": "NONE" },
      { "severityLevel": "CUSTOM_ALERT", "delayInMinutes": 0, "tagFilterIncludeMode": "NONE" }
    ],
    "eventFilters": [],
    "managementZone": "<management-zone-object-id>"
  }
}
```

3.4.9. Host CPU metric event (optional companion path)

Host CPU problems are **not Spark-specific** unless scoped. Use a **neutral event title** and an **entity filter** so the detector evaluates only hosts in the integration boundary (management zone or named hosts). Different business units may use different thresholds — create **separate metric events** per scope rather than one tenant-wide rule.

```
{
  "schemaId": "builtin:anomaly-detection.metric-events",
  "scope": "environment",
  "value": {
    "enabled": true,
    "summary": "Spark Lab - Host CPU above 80%",
    "queryDefinition": {
      "type": "METRIC_KEY",
      "metricKey": "builtin:host.cpu.usage",
      "aggregation": "AVG",
      "entityFilter": {
        "conditions": [
          {
            "type": "HOST_GROUP_NAME",
            "operator": "EQUALS",
            "value": "<integration-host-group>"
          }
        ]
      },
      "dimensionFilter": []
    },
    "eventEntityDimensionKey": "dt.entity.host",
    "modelProperties": {
      "type": "STATIC_THRESHOLD",
      "threshold": 80.0,
      "alertOnNoData": false,
      "alertCondition": "ABOVE",
      "violatingSamples": 1,
      "samples": 3,
      "dealertingSamples": 3
    },
    "eventTemplate": {
      "title": "Host ({dims:dt.entity.host}) CPU above 80%",
      "description": "Host CPU usage in management zone <name> exceeded 80%.",
      "eventType": "CUSTOM_ALERT",
      "davisMerge": false,
      "metadata": []
    }
  }
}
```

Scope the detector with an **entity filter**: **HOST_GROUP_NAME** (when OneAgent assigns a host group), **MANAGEMENT_ZONE** (numeric zone id), or individual **HOST_NAME** values. Multiple conditions in one filter use **AND** logic — use a single host group or zone rather than listing every host unless each host has its own metric event.

When the integration uses a **management zone**, a **MANAGEMENT_ZONE** filter requires the zone's **numeric id** (not the Settings API `objectId`). Template: <observability/dynatrace/tenants/<tenant>/integrations/spark-cpu-metric-event.json.j2>.

3.4.10. ServiceNow (Step 2)

No direct ServiceNow configuration. Ensure SGC scheduled imports cover `CLOUD_APPLICATION_INSTANCE` and `PROCESS_GROUP` entity types that Davis may select as `dt.source_entity`.

3.5. Step 3 — Dynatrace sends the problem to ServiceNow

3.5.1. Webhook endpoint

```
POST https://<instance>.service-now.com/api/sn_em_connector/em/inbound_event?source=SG0-Dynatrace
Authorization: Basic <base64(user:password)>
Content-Type: application/json
```

Use **ProblemDetailsJSON** (v1) in the payload template for SGC — not **ProblemDetailsJSONv2** (legacy templates may return HTTP 500 on some instances).

3.5.2. Problem notification payload template

The template in [Problem notification payload specification](#) is substituted by Dynatrace at send time. **ProblemDetailsJSON.rankedEvents[]** carries the OpenPipeline Davis properties from [OpenPipeline specification](#) (`event.description`, entity ids, severity). Those ranked events are the **problem specification** that flows into ServiceNow `em_event.description` (via **ProblemDetailsText**) and into `em_event.additional_info`.

Specification 2.6.2-1. Problem notification payload (SGC)

```
{
  "ConnectionId": "<sgc-connection-sys-id>",
  "ProblemID": "{ProblemID}",
  "ProblemTitle": "{ProblemTitle}",
  "ProblemSeverity": "{ProblemSeverity}",
  "ProblemImpact": "{ProblemImpact}",
  "ProblemDetailsJSON": {ProblemDetailsJSON},
  "ProblemDetailsText": "{ProblemDetailsText}",
  "ImpactedEntities": {ImpactedEntities},
  "ImpactedEntity": "{ImpactedEntity}",
  "ProblemURL": "{ProblemURL}",
  "State": "{State}",
  "Tags": "{Tags}",
  "PID": "{PID}"
}
```

Enable **notifyClosedProblems** so RESOLVED posts update the same `em_event` via `message_key`.

3.5.3. Webhook field mapping

Webhook field	Dynatrace source	ServiceNow target
ProblemID	Stable problem identifier	em_event.message_key (OPEN and RESOLVED deduplication)
ProblemTitle	event.name / Davis title	em_event.message; contributes to description
ProblemSeverity	ERROR, CUSTOM_ALERT, etc.	em_event.severity (numeric mapping per SGC rules)
ImpactedEntities[].entityId	CLOUD_APPLICATION_INSTANCE-..., PROCESS_GROUP-..., HOST-...	sys_object_source lookup → em_event.cmdb_ci
ProblemDetailsJSON.rankEvents[]	Ranked events with entityId, eventType	Fallback CI binding parse path
ProblemDetailsText	Davis narrative including event.description	em_event.description (primary operator-facing text)
Tags	Problem tags (Project, Environment, ...)	em_event.additional_info; not CSDM join keys
State	OPEN or RESOLVED	Create vs close/update event
ProblemURL	Deep link to Dynatrace problem	em_event.additional_info

3.5.4. Complete example webhook body

After Dynatrace substitutes placeholders for an ERROR log problem:

```
{
  "ConnectionId": "a1b2c3d4e5f678901234567890abcdef0",
  "ProblemID": "-9876543210987654321",
  "ProblemTitle": "Application log ERROR",
  "ProblemSeverity": "ERROR",
  "ProblemImpact": "INFRASTRUCTURE",
  "ProblemDetailsJSON": {
    "displayName": "Application log ERROR",
    "problemId": "-9876543210987654321",
    "rankedEvents": [
      {
        "entityId": "CLOUD_APPLICATION_INSTANCE-A1B2C3D4E5F67890",
        "entityName": "myapp-worker-7d4f9b-xk2lm",
        "eventType": "ERROR_EVENT",
        "severityLevel": "ERROR",
        "title": "Application log ERROR",
        "eventDescription": "Application ERROR on /var/log/myapp/myapp-worker-7d4f9b-xk2lm/application.log: 2026-06-27 14:53:46 ERROR Worker:264 - All masters are unresponsive! Giving up."
      }
    ],
    "startTime": 1719502426000,
    "endTime": -1
  },
  "ProblemDetailsText": "Application ERROR on /var/log/myapp/myapp-worker-7d4f9b-xk2lm/application.log:\n2026-06-27 14:53:46 ERROR Worker:264 - All masters are unresponsive! Giving up.\n\n1 impacted entity: myapp-worker-7d4f9b-xk2lm",
  "ImpactedEntities": [
    {
      "entityId": "CLOUD_APPLICATION_INSTANCE-A1B2C3D4E5F67890",
      "name": "myapp-worker-7d4f9b-xk2lm",
      "type": "CLOUD_APPLICATION_INSTANCE"
    }
  ]
}
```

```

    }
  ],
  "ImpactedEntity": "myapp-worker-7d4f9b-xk2lm",
  "ProblemURL": "https://<tenant>.apps.dynatrace.com/ui/problems/<problem-id>",
  "State": "OPEN",
  "Tags": "Environment:production,Project:myapp",
  "PID": "-9876543210987654321"
}

```

3.5.5. Ensuring log fields reach ServiceNow

Layer	Configuration
OpenPipeline Davis processor	Set event.description to include {content} , {loglevel} , and {log.source} (Step 2); set event.name to include {k8s.pod_name} for per-pod correlation (OpenPipeline specification)
Dynatrace problem notification	When OpenPipeline Davis properties are present, Dynatrace expands them into ProblemDetailsJSON.rankedEvents[] (each ranked event includes eventDescription , entityId , title). Dynatrace then assembles ProblemDetailsText from those ranked events (including event.description text) before the webhook is sent. See Problem notification payload specification .
SGC field mapping (sa_event_field_mapping in sn_dynatrace_integ)	Map ProblemDetailsText → em_event.description ; map ImpactedEntity or parsed pod name → em_event.resource
em_event business rule (Spark pod bind)	When path or node encodes pod name, set cmdb_ci , resource , message_key prefix K8sLog-{pod}- , and message Event Log WARN or Event Log ERROR
em_event.additional_info	Persist full ProblemDetailsJSON for Event Management rules and incident scripts
em_alert.description	Event Management alert rule copies em_event.description ; cmdb_ci copied from event



Kubernetes labels (servicenow.io/application-service-identifier) are **not** in the webhook payload. They remain on CMDB pod CIs for Step 5 incident correlation.

3.5.6. ServiceNow (Step 3)

The **Observability engineer** (ServiceNow administrator with Event Management and SGC access) must verify the push connector — not the monitored application developer:

1. Confirm the Store app **Service Graph Connector for Observability – Dynatrace** (**sn_dynatrace_integ**) is installed and active.
2. In **Event Management** → **Integrations** → **Push Connectors** (or **Connections** → **Observability** → **Dynatrace** in newer releases), locate the connector whose inbound URL ends with:

```
POST /api/sn_em_connector/em/inbound_event?source=SGO-Dynatrace
```

3. Verify **Basic authentication** credentials match the Dynatrace problem notification profile

(brooks-lab: provisioned via Ansible `ensure_sn_dt_webhook_auth.yml`).

4. In the scoped app, open **Event field mappings** (`sa_event_field_mapping / $sa_event_map`) and confirm: **ProblemDetailsText** → **description**; **ProblemTitle** → **message**; **ProblemID** → **message_key**; severity mapping for ERROR and WARN.
5. Ensure EM **event rules** mark inbound SGO-Dynatrace events **Ready** when `cmdb_ci` is populated (see [Event Management rules](#)).
6. Ensure EM **alert rules** create alerts from Ready events and copy `cmdb_ci`, `description`, `resource`, and `node` from the event.

No **ServiceNow** configuration is required on the monitored application itself. No **Dynatrace** changes are required at this step beyond the problem notification profile deployed in [Step 2 — Dynatrace](#).

3.6. Step 4 — ServiceNow receives the event and binds infrastructure CI

3.6.1. Event listener

Setting	Value
Source	SGO-Dynatrace (URL query parameter on <code>inbound_event</code>)
Scoped app	<code>sn_dynatrace_integ</code>
CI lookup	<code>sys_object_source</code> where <code>name=SGO-Dynatrace</code> and <code>id=<entityId from ImpactedEntities></code>

3.6.2. `em_event` field mapping

<code>em_event</code> field	Population
<code>source</code>	SGO-Dynatrace (from URL)
<code>message_key</code>	<code>ProblemID</code>
<code>description</code>	<code>ProblemDetailsText</code> — must include full log line from <code>OpenPipeline event.description</code>
<code>message</code>	<code>ProblemTitle</code>
<code>severity</code>	Mapped from <code>ProblemSeverity</code> (ERROR → 2 Major typical; WARN → 3 Minor if applicable)
<code>type</code>	ERROR or derived from <code>ProblemDetailsJSON</code> ranked event type
<code>node</code>	<code>ImpactedEntity</code> display name (pod name when <code>CLOUD_APPLICATION_INSTANCE</code>)
<code>resource</code>	Log file path or <code>workload.pod_name</code> parsed from <code>ProblemDetailsText</code> or <code>additional_info</code>

em_event field	Population
cmdb_ci	sys_object_source.target_sys_id for entityId in ImpactedEntities
cmdb_ci_type	cmdb_ci_kubernetes_pod when entity is CLOUD_APPLICATION_INSTANCE; cmdb_ci_appl for PROCESS_GROUP; cmdb_ci_linux_server for HOST fallback
additional_info	JSON: ProblemURL, Tags, ImpactedEntities, full ProblemDetailsJSON, PID
time_of_event	Problem start time from ProblemDetailsJSON
correlation_id	PID when mapped

3.6.3. CI binding procedure

1. Listener receives POST at inbound_event?source=SGO-Dynatrace.
2. Parser extracts entityId from ImpactedEntities[0] or ProblemDetailsJSON.rankedEvents[0].
3. Query sys_object_source: name=SGO-Dynatrace, id=<entityId>.
4. Set em_event.cmdb_ci and em_event.cmdb_ci_type from the matched row.

3.6.4. Entity type to CMDB class

Dynatrace entityId prefix	CMDB class	When to expect
CLOUD_APPLICATION_INSTANCE	cmdb_ci_kubernetes_pod	OpenPipeline resolved pod entity (preferred for K8s log paths)
PROCESS_GROUP	cmdb_ci_appl	Log tied to JVM process when pod lookup unavailable
HOST	cmdb_ci_linux_server	Last resort — file tailed on node without pod enrichment

Prerequisites for pod binding:

- SGC Kubernetes Pod import includes pods in the application namespace.
- Pod name in Dynatrace matches Kubernetes metadata.name.
- IRE merge aligns KVA-discovered pod CIs with SGC-imported entities when both exist.
- sys_object_source row exists before the webhook arrives (account for import schedule lag).

3.6.5. Description and alert text

Field	Required content
em_event.description	Line 1: log level and message body from content; Line 2: log.source path; optional ProblemURL appended
em_alert.description	Same as em_event.description (copied by alert rule or inherited on alert create)
em_alert.resource	Pod name or full log file path
em_alert.cmdb_ci	Same infrastructure CI as em_event.cmdb_ci

3.6.6. SGC and Event Management configuration

Component	Specification
sa_event_field_mapping	ProblemDetailsText → description; ProblemTitle → message; ProblemID → message_key; severity mapping for ERROR and WARN
EM event rule	Mark events Ready when severity and cmdb_ci are populated; optional script normalizes description from additional_info if connector mapping is insufficient
EM alert rule	Create alert from Ready events; copy cmdb_ci and description from event; require cmdb_ci for incident automation
ACL on cmdb_key_value	Roles that run incident scripts must read tags on pod CIs (see install documentation for tag ACL requirements)

3.6.7. Application and Dynatrace (Step 4)

At Step 4 the **Observability engineer** typically performs **verification**, not new Dynatrace or application configuration:

1. After a test problem fires, confirm `em_event.source=SGO-Dynatrace` and `em_event.cmdb_ci` is populated (pod CI when OpenPipeline entity lookup succeeded).
2. If `cmdb_ci` is a Linux server instead of a pod, return to [Step 2 — Dynatrace](#) and verify `k8s.pod_name` extraction and `lookupEntity(CLOUD_APPLICATION_INSTANCE, k8s.pod_name)` in the OpenPipeline pipeline ([OpenPipeline specification](#)).
3. Confirm `sys_object_source` contains a row with `name=SGO-Dynatrace` and `id=<entityId>` from the webhook **before** the event arrives (SGC scheduled import lag is a common root cause of empty CI).
4. Deploy or verify the **K8s log pod bind** business rules ([K8s log pod bind business rules](#)) as a fallback when the webhook still reports a HOST entity but the log path encodes a pod name.

No **ServiceNow-side changes** are required on the monitored **application** (the Java/Spark/Python workload). No **additional Dynatrace** configuration is required at Step 4 if Steps [Step 2 — Dynatrace](#) and Step 3 are correct.

ServiceNow cannot bind to `cmdb_ci_kubernetes_pod` unless [Step 2 — Dynatrace](#) sets `dt.source_entity` to the pod entity and Step 0 imported that entity into `sys_object_source`.

3.7. Step 5 — Create incident bound to Application Service

3.7.1. Severity policy

Environment	Automatic incident when
Non-production / test	ERROR and WARN (simplifies validation when ERROR signals are hard to trigger)
Production	ERROR only; WARN remains at alert level or creates a lower-priority task per ITSM policy

Incidents must set **incident.cmdb_ci** to **cmdb_ci_service_discovered** (Application Service), not the infrastructure CI on the alert.

3.7.2. Layer model

Layer	Table	cmdb_ci class
Event	em_event	Infrastructure: cmdb_ci_kubernetes_pod, cmdb_ci_appl, or cmdb_ci_linux_server
Alert	em_alert	Same infrastructure CI as the source event
Incident	incident	cmdb_ci_service_discovered (Application Service)

3.7.3. Incident correlation

Incident correlation defines how enriched alerts map to **incident.cmdb_ci**. The business rule **em-alert-create-k8s-log-incident** tries the **Client-side** log path first, then the **Service-side** Contains traversal.

The comparison table summarizes alert CI versus incident CI for each pattern. Full algorithms follow in the adjacent subsections.

Pattern	Alert CI	incident.cmdb_ci resolution
Service-side	cmdb_ci_kubernetes_pod (preferred) or HOST fallback	Spark Service-Side Log Pattern — Contains traversal from pod CI
Client-side	HOST or empty (acceptable)	Spark Client-Side Log Pattern — /logs/spark-client/ → Spark Client AS

3.7.3.1. Service-side

The **Spark Service-Side Log Pattern** resolves Application Services by traversing **Contains::Contained by** from the alert's pod CI. A materialized edge from **cmdb_ci_service_discovered** (parent) to **cmdb_ci_kubernetes_pod** (child) **must** exist before the problem fires.



Figure 3. Incident correlation — Spark Service-Side Log Pattern

1. **K8sLogPodCiBind** sets **em_event.cmdb_ci** and **em_alert.cmdb_ci** to the **pod CI** whose **name** matches

the log path segment (Step 4).

2. Query `cmdb_rel_ci` where `child` = pod CI `sys_id` and `type` = **Contains::Contained by**.
3. Parent row **must** be `cmdb_ci_service_discovered`.
4. Set `incident.cmdb_ci` = parent Application Service `sys_id`.

Out of scope: host CI binding without pod rebind, process group CI, tag-only lookup without a Contains edge.

3.7.3.2. Client-side

The **Spark Client-Side Log Pattern** maps log paths to a single logical Application Service. Client-mode driver logs **do not** correspond to a `cmdb_ci_kubernetes_pod`; incident automation **must not** depend on alert CI or Contains edges.




Figure 4. Incident correlation — Spark Client-Side Log Pattern

1. Scan `em_alert.description`, `resource`, and `node` for the substring `/logs/spark-client/`.
2. Look up Application Service **Spark Client** by `name` (see [CSDM identifier lookup](#)).
3. Set `incident.cmdb_ci` to that Application Service `sys_id`.
4. `K8sLogPodCiBind` skips path segment `spark-client` when parsing pod names.

3.7.4. CSDM identifier lookup

CSDM **identifier** in `*.csdm.yaml` is the logical machine key embedded in log paths and optional Dynatrace tags. The table below shows which CMDB fields are reliable for incident lookup on `optimizincdemo1`.

Purpose	Mechanism	Queryable field on <code>cmdb_ci_service_discovered</code>
Manual / logical AS documentation	CSDM identifier : <code>spark-client</code> in log path contract	identifier column — often not populated on <code>optimizincdemo1</code>
Display / incident lookup	CSDM name (e.g. <code>Spark Client</code>)	name — reliable for <code>resolveFromSparkClientLogPath()</code>
Optional process tag	<code>DT_TAGS=service-now.io/application-service-identifier=spark-client</code>	Process-group problems only; not required for log-path incidents

Spark Client AS **must** have `service_status=operational` (promoted by `csdm/deploy.yaml` → `ensure_host_as_operational.yaml`).

```
cd ansible
ansible-playbook -i inventory.yml playbooks/service-now/incident/verify_spark_client_as.yml \
  -e @../vars/secrets.yaml
```


3.7.5. Known limitation — Davis problem bundling

Dynatrace Davis **bundles** multiple log-derived `ERROR_EVENT` rows that share a root-cause analysis into **one problem** when they share the same `dt.source_entity` (historically the OneAgent **HOST** for file-tailed logs). When client-mode and Service-side log lines fire within the same Davis correlation window on the same host, a **single** `em_alert.description` can contain **both**:

- Application log WARN on spark-client-... with path `/mnt/spark/logs/spark-client/...`
- Application log WARN on spark-master-0 with path `/mnt/spark/logs/spark-master-0/...`

Remediation (lab):

- OpenPipeline `resolve-spark-client-entity` sets client `dt.source_entity` via `lookupEntity(CUSTOM_DEVICE, "Spark Client")` (logical device, not HOST). See `Davis_Bundling_Issue.md`.
- ServiceNow `ResolveApplicationService.applySparkClientAlertBinding()` binds alert/event CI to Application Service **Spark Client** when `/logs/spark-client/` is present and **skips** pod rebind.

Historical impact (pre-remediation): Client-side and Service-side validation in the same test window produced **misleading alert CI** (`spark-master-0`) and incidents on the host. Split test runs or wait for Davis timeout (lab: `15m dt.davis.event_timeout`) remain useful as isolation controls.

3.7.6. ServiceNow incident automation specification

Component	Specification
Script Include: <code>ResolveApplicationService</code>	Service-Side: <code>cmdb_rel_ci</code> Contains::Contained by from pod CI to Application Service. Spark client: <code>resolveFromSparkClientLogPath()</code> when alert text contains <code>/logs/spark-client/</code> → Application Service Spark Client
Script Include: <code>K8sLogPodCiBind</code>	Rebind <code>em_event</code> / <code>em_alert</code> from initial SGO HOST binding to <code>cmdb_ci_kubernetes_pod</code> when log path segment matches a discovered pod name; skip segment <code>spark-client</code> ; propagate binding to alerts sharing <code>message_key</code>
EM event rule	Mark Spark log events Ready when <code>source=SG0-Dynatrace</code> , severity matches policy, and <code>cmdb_ci</code> is set
EM alert rule	Create alert from Ready events; copy <code>cmdb_ci</code> , <code>description</code> , <code>resource</code> , <code>node</code> from event
Business rule: <code>em-event-bind-k8s-log-pod-ci</code>	Before insert/update on <code>em_event</code> ; order 5000 ; enriches <code>cmdb_ci</code> , <code>resource</code> , <code>message_key</code> via <code>K8sLogPodCiBind</code> (must run before EM Ready/alert rules persist the row)
Business rule: <code>em-event-propagate-k8s-log-pod-ci</code>	After insert/update on <code>em_event</code> ; order 5005 ; copies pod binding from event to related <code>em_alert</code> rows sharing <code>message_key</code> (runs after the event row is saved)
Business rule: <code>em-alert-bind-k8s-log-pod-ci</code>	Before insert/update on <code>em_alert</code> ; order 5010 ; same pod/path enrichment when alert is created or updated directly (covers SGC webhook path that skips a separate bind on the event)

Component	Specification
Business rule: <code>em-alert-create-k8s-log-incident</code>	After insert/update on <code>em_alert</code> ; order 5020 ; tries spark-client path first, else Service-Side Contains; creates or correlates incident, sets <code>em_alert.incident</code>
Severity filter	Non-prod: <code>em_alert.severity</code> $\leq$ 3 (ERROR and WARN); Production: <code>em_alert.severity</code> = 2 (ERROR only)
Duplicate suppression	Correlate open incidents by Application Service + short description prefix <code>Event Log</code> ; use <code>ProblemID</code> / <code>message_key</code> for event lifecycle

3.7.7. Business rule order and rationale

Business rules on `em_event` and `em_alert` run in **order** within the same **when** phase (**before** or **after**). Lower order runs first.

Order	Rule	Why this position
5000	<code>em-event-bind-k8s-log-pod-ci</code> (before)	Parse log path and set <code>cmdb_ci</code> / <code>resource</code> / <code>message_key</code> before the row is written. EM event rules that mark the event Ready and copy fields to alerts depend on these values being present on insert.
5005	<code>em-event-propagate-k8s-log-pod-ci</code> (after)	After the event is saved, push binding to sibling <code>em_alert</code> rows. Must be after (not before) because it updates other records.
5010	<code>em-alert-bind-k8s-log-pod-ci</code> (before)	Same enrichment for alerts created or updated by the SGC webhook / EM alert rule without going through a separate event bind cycle.
5020	<code>em-alert-create-k8s-log-incident</code> (before)	Needs enriched <code>resource</code> and log-path context from 5000/5010 in the same transaction. Creates the incident and sets <code>em_alert.incident</code> before the alert row is committed. Using after + <code>current.update()</code> was unreliable for REST-driven updates.

3.7.8. Alert to incident linkage (`em_alert.incident`)

On `optimizincdemo1`, Event Management links an alert to its incident through `em_alert.incident` (reference to the incident record). That field populates the incident **Alerts** tab. Some instances expose the same reference as `em_alert.task`; this build uses **incident only** — do not reference `task`.

Creating an incident with `GlideRecord('incident').insert()` alone does **not** set this link — automation must assign `current.incident = <incident_sys_id>` on the alert in the **before** business rule.

Incident short description format: `Event Log WARN` or `Event Log ERROR` (parsed from alert description).

3.7.9. Problem and alert lifecycle (OPEN / RESOLVED)

Spark/Log4j does **not** emit explicit “clear” or “recovery” log lines for this use case. Closure is driven by **Dynatrace problem lifecycle**, not application log content:

1. OneAgent tails application log files on the **host node** (see [Log emission](#)).

2. OpenPipeline Davis turns qualifying **ERROR** / **WARN** lines into Davis events; repeated lines merge into one **problem** (`dt.davis.event_timeout: 15m`).
3. When matching log activity stops, Dynatrace **auto-closes** the problem after the timeout.
4. The brooks-lab problem notification has `notifyClosedProblems: true`, so Dynatrace POSTs a **RESOLVED** webhook with the same **ProblemID** as the OPEN post.
5. SGO maps `message_key = ProblemID`, so ServiceNow **updates the same em_event** (does not create a duplicate).
6. Event Management then transitions the linked **alert** (and optionally the **incident**) according to configured EM rules when the event state/resolution changes.

Verify RESOLVED handling: `em_event` where `source=SGO-Dynatrace` and `resolution_code` is set after problems clear; linked `em_alert.state` should move to **Closed** when EM auto-close rules are configured.

3.7.10. OpenPipeline and custom-source attributes (generic Log4j pattern)

Attributes added by automation (not Spark-specific semantics):

Attribute	Source	Purpose
<code>k8s.log.path</code>	Custom log source enrichment	Full file path of the log line
<code>k8s.pod_name</code>	Path segment after <code>/logs/</code>	Must equal <code>cmdb_ci_kubernetes_pod.name</code>
<code>event.name</code>	OpenPipeline Davis	Application log {loglevel} on {k8s.pod_name}
<code>event.description</code>	OpenPipeline Davis	Application {loglevel} on {k8s.log.path}: {content}

Pod **Kubernetes labels** (`servicenow.io/application-service-identifier`, etc.) are the CMDB/service-mapping contract — not Dynatrace tags. See [pod labels](#).



resolve-k8s-pod-entity (OpenPipeline `lookupEntity` for pod Smartscape) is **disabled** so problems stay on **HOST** entities inside the Spark Observability management zone. ServiceNow `K8sLogPodCiBind` rebinds alerts to the **pod CI** from the log path; `ResolveApplicationService` traverses the **Contains::Contained** by edge to the Application Service (Service-Side pattern only).

3.7.11. Event Management rules (manual configuration)

Configure under **Event Management** → **Rules**:

Rule	Condition and action
Event rule (K8s application logs)	When: <code>em_event</code> insert/update; Filter: <code>source=SGO-Dynatrace</code> <code>severity<=3</code> <code>cmdb_ciISNOTEMPTY</code> ; Action: set state to Ready (or use existing Ready rule scoped to SGO-Dynatrace)

Rule	Condition and action
Alert rule (K8s application logs)	When: Ready event; Filter: <code>source=SGO-Dynatrace</code> ; Action: create/update alert; copy <code>cmdb_ci</code> , <code>description</code> , <code>resource</code> , <code>node</code> , <code>message_key</code> from event

Scope Spark-specific business rules with `source=SGO-Dynatrace` so legacy `source=Dynatrace` events are unaffected during connector migration (see [Legacy and SGO-Dynatrace coexistence](#)).

3.7.12. Script Includes

Deploy from `servicenow/integrations/incident/`. Global scope, **Active** = true. Ansible: `ansible/playbooks/servicenow/incident/deploy.yml`.

Each Script Include below summarizes **purpose**, **inputs**, **outputs** / **side effects**, and how it fits the business rules. Source of truth is the `.si.js` file in the repository.

3.7.12.1. ResolveApplicationService

Purpose: Resolve `cmdb_ci_service_discovered` (Application Service) `sys_id` for incident binding.

Methods:

- `resolveFromInfrastructureCi(ciSysId)` — **Spark Service-Side Log Pattern**. Queries `cmdb_rel_ci` for **Contains::Contained by** where `child` = pod CI and `parent` is an Application Service. Returns parent `sys_id` or `null`.
- `resolveFromSparkClientLogPath(text)` — **Spark Client-Side Log Pattern**. When `text` contains `/logs/spark-client/`, looks up Application Service **Spark Client** by `identifier=spark-client`, then fallback `name=Spark Client`.

Inputs: Workload CI `sys_id` (Service-side) or alert description/resource text (Client-side).

Outputs: Application Service `sys_id` for `incident.cmdb_ci`; no CMDB writes.

Side effects: None (read-only GlideRecord queries).

Source: `servicenow/integrations/incident/ResolveApplicationService.si.js`

3.7.12.2. K8sLogPodCiBind

Purpose: Rebind `em_event` and `em_alert` from HOST or empty CI to `cmdb_ci_kubernetes_pod` when the log path directory segment matches a discovered pod name (OpenPipeline often leaves problems on HOST).

Methods:

- `resolvePodName(description, resource, node)` — Scans concatenated fields for `/.../logs/<segment>/`; skips `spark-client` and wildcards; verifies pod exists in CMDB.
- `applyPodBinding(gr)` — Sets `cmdb_ci`, `node`, `resource`, `message_key` prefix `K8sLog-{pod}-`, and `message` to Event Log **WARN/ERROR** when generic. Returns `true` on success. Caller must `update()`.

- `propagateEventToAlerts(eventGr)` — Copies pod binding from event to related alerts sharing `message_key`.

Inputs: `em_event` or `em_alert` GlideRecord.

Outputs: Mutated fields on current row; optional `em_alert.update()` on related rows.

Source: `servicenow/integrations/incident/K8sLogPodCiBind.si.js`

3.7.13. Business rules (prescriptive configuration)

Create under **System Definition** → **Business Rules**. All three run **After** insert and update. Filter conditions must include `source=SGO-Dynatrace` so legacy connectors are not modified.

Name	Table	Order	Filter condition
<code>em-event-bind-k8s-log-pod-ci</code>	<code>em_event</code>	5000	<code>source=SGO-Dynatrace</code>
<code>em-alert-bind-k8s-log-pod-ci</code>	<code>em_alert</code>	5010	<code>source=SGO-Dynatrace</code>
<code>em-alert-create-k8s-log-incident</code>	<code>em_alert</code>	5100	<code>source=SGO-Dynatrace^severity<=3</code>

3.7.13.1. em-event-bind-k8s-log-pod-ci

```
(function executeRule(current, previous /*null on insert*/) {
  var binder = new K8sLogPodCiBind();
  if (!binder.applyPodBinding(current)) {
    return;
  }

  current.update();
  binder.propagateEventToAlerts(current);
})(current, previous);
```

3.7.13.2. em-alert-bind-k8s-log-pod-ci

```
(function executeRule(current, previous /*null on insert*/) {
  var binder = new K8sLogPodCiBind();
  if (!binder.applyPodBinding(current)) {
    return;
  }

  current.update();
})(current, previous);
```

3.7.13.3. em-alert-create-k8s-log-incident

Creates or correlates an incident on the Application Service, then sets `current.task` so the alert appears on the incident **Alerts** tab.

```
(function executeRule(current, previous /*null on insert*/) {
  if (current.source.toString() !== 'SGO-Dynatrace') {
    return;
  }
}
```

```

}

if (!current.task.nil()) {
    return;
}

var severity = parseInt(current.severity, 10);
if (isNaN(severity) || severity > 3) {
    return;
}

var desc = current.description.toString();
if (
    desc.indexOf('Spark ') === -1 &&
    desc.indexOf('spark-app.log') === -1 &&
    desc.indexOf('/mnt/spark/logs/') === -1
) {
    return;
}

if (current.cmdb_ci.nil()) {
    return;
}

var resolver = new ResolveApplicationService();
var asSysId = resolver.resolveFromInfrastructureCi(current.cmdb_ci.toString());
if (!asSysId) {
    return;
}

var levelMatch = desc.match(/\b(ERROR|WARN)\b/);
var logLevel = levelMatch
    ? levelMatch[1]
    : severity <= 2
    ? 'ERROR'
    : 'WARN';
var shortDesc = 'Event Log ' + logLevel;
var shortDescPrefix = 'Event Log';

function linkAlertToIncident(incSysId) {
    var alertGr = new GlideRecord('em_alert');
    if (!alertGr.get(current.sys_id)) {
        return;
    }
    alertGr.task = incSysId;
    alertGr.setWorkflow(false);
    alertGr.update();
}

var backfillInc = new GlideRecord('incident');
backfillInc.addQuery('active', true);
backfillInc.addQuery('short_description', 'STARTSWITH', shortDescPrefix);
backfillInc.addEncodedQuery('cmdb_ciIEMPTY');
backfillInc.orderBy('sys_created_on');
backfillInc.setLimit(1);
backfillInc.query();
if (backfillInc.next()) {
    backfillInc.cmdb_ci = asSysId;
    backfillInc.short_description = shortDesc;
    backfillInc.work_notes =
        'Set Configuration item from alert ' +
        current.number +
        ' (Application Service resolved from pod CI).';
    backfillInc.update();
    linkAlertToIncident(backfillInc.sys_id);
    return;
}

```

```

var openInc = new GlideRecord('incident');
openInc.addQuery('cmdb_ci', asSysId);
openInc.addQuery('active', true);
openInc.addQuery('short_description', 'STARTSWITH', shortDescPrefix);
openInc.setLimit(1);
openInc.query();
if (openInc.next()) {
    openInc.work_notes =
        'Correlated alert ' +
        current.number +
        ': ' +
        current.description.toString().substring(0, 500);
    openInc.update();
    linkAlertToIncident(openInc.sys_id);
    return;
}

var inc = new GlideRecord('incident');
inc.initialize();
inc.cmdb_ci = asSysId;
inc.short_description = shortDesc;
inc.description = current.description;
inc.severity = current.severity;
inc.impact = 2;
inc.urgency = 2;
inc.caller_id = gs.getUserID();
inc.comments = 'Auto-created from Event Management alert ' + current.number;
inc.work_notes = 'Source alert: ' + current.number;
var incSysId = inc.insert();
linkAlertToIncident(incSysId);
})(current, previous);

```

3.7.14. Ansible deployment

```

cd ansible
ansible-playbook -i inventory.yml playbooks/servicenow/incident/deploy.yml \
-e @../vars/secrets.yaml

```

Playbook path: [ansible/playbooks/servicenow/incident/deploy.yml](#). Source files: [servicenow/integrations/incident/*.js](#).

3.7.15. Example resolution (Kubernetes worker pod)

Step	Record / relationship
Problem	Davis problem; impacted entity CLOUD_APPLICATION_INSTANCE-A1B2C3D4E5F67890
Event	em_event; cmdb_ci → cmdb_ci_kubernetes_pod myapp-worker-7d4f9b-xk2lm
Alert	em_alert; same pod CI; description contains full ERROR log line
Tag on pod CI	cmdb_key_value: servicenow.io/application-service-identifier=myapp-worker-zone-a
SM edge	Application Service myapp-worker-zone-a Contains pod CI
Incident	incident.cmdb_ci → Application Service myapp-worker-zone-a (via Service-Side Contains traversal)

3.7.16. Application and Dynatrace (Step 5)

No configuration on the application or Dynatrace side. Dynatrace problem tags (Project, Environment) do not replace CSDM tags on pod CIs for incident binding.

Appendix A: Legacy and SGO-Dynatrace coexistence

During migration from a legacy Event Management push connector (`source=Dynatrace` on the inbound URL) to Service Graph Connector (`source=SGO-Dynatrace`), **both paths may be active**. Use `em_event.source` and `em_alert.source` to scope automation — never assume a single connector.

A.1. Separation by source

Layer	Legacy (<code>source=Dynatrace</code>)	SGC (<code>source=SGO-Dynatrace</code>)
Inbound URL	<code>/inbound_event?source=dynatrace&sys_id=<connector-sys-id></code>	<code>/inbound_event?source=SGO-Dynatrace</code>
CI binding	Legacy connector field mapping + <code>sys_object_source</code> (may differ from SGC import)	SGC <code>ImpactedEntities</code> → <code>sys_object_source</code> (name SGO-Dynatrace)
Spark log business rules	Must not run — filter <code>source=SGO-Dynatrace</code> on every BR	Pod rebind, incident create, Application Service resolution
Problem notification	Pre-existing tenant profiles (for example Demo integrations)	Dedicated profile scoped to integration management zone

A.2. Scoping rules and profiles

- **Dynatrace alerting profiles:** Scope each profile to a **management zone** (or tag filter) so only in-scope problems forward to ServiceNow. A profile attached to a legacy webhook should not receive problems from hosts outside its intended boundary.
- **ServiceNow EM rules:** Add `source=SGO-Dynatrace` (or `source=Dynatrace`) to event and alert rule filter conditions so each rule set applies only to its connector.
- **Business rules:** All Spark log automation in this document filters `source=SGO-Dynatrace`. Legacy CPU or infrastructure alerts continue through legacy rules until that connector is retired.

A.3. Host CPU alerts (not Spark-specific)

Metric events on `builtin:host.cpu.usage` fire per **host**, not per application. Even on the SGC path, a CPU problem binds `em_alert.cmdb_ci` to the breaching **host CI** — that is correct behavior.

To avoid misleading titles and out-of-scope hosts:

1. Use a **neutral title**: `Host ({dims:dt.entity.host}) CPU above 80%` (not application-branded text).
2. Add an **entity filter** on the metric event: **management zone** (preferred) or explicit **HOST_NAME** list for integration hosts.

3. Create **separate metric events** per organization or environment when thresholds differ (for example 80% in lab, 90% in production).

Template: `observability/dynatrace/tenants/<tenant>/integrations/spark-cpu-metric-event.json.j2`. Deploy via `ansible/playbooks/servicenow/sgc/sources/dynatrace/tasks/apply_dt_anomaly_detectors.yml`.

A.4. When to retire legacy connector

After SGC cutover is verified (events arrive with `source=SGO-Dynatrace`, pod CI binding works, incidents link on **Alerts** tab):

1. Disable legacy Dynatrace problem notification profiles pointing at `source=dynatrace`.
 2. Leave legacy EM rules in place until no new `source=Dynatrace` events arrive.
 3. Remove legacy rules once the connector is decommissioned.
-